

CS 33:
Introduction to Computer
Organization

Week 3

Agenda

- x86 Organization
- x86 Assembly

x86 Organization

- The basic abstraction of memory that is taught in CS 31 and CS 32 is that data is stored in memory.
- For example, if you have a 32-bit addressable space, then the addresses that are in memory range from 0x00000000 to 0xFFFFFFFF($2^{31}-1$).
- It's not like you'll ever see a variable that you can't dereference in C. Therefore variables must always be stored in memory, right?

x86 Organization

- The variables will be stored in memory, which is a physical construct (RAM).
- However, RAM is too slow to keep up with the demands of a processor.
- Accessing RAM takes approximately 200 times the amount of time as it takes to execute a standard instruction.

x86 Organization

- Since nearly everything involves doing some operation on a variable, we need some way of accessing memory at a speed that is comparable to the speed that it takes to execute the average instruction.
- We need caches, we need virtual memory, but for now, we'll focus on “registers”.

x86 Organization

- Registers are extremely small physical containers that each store a number of bits.
- A 64-bit addressable machine will have registers that are 64-bits.
- In such a case, each register holds 8 bytes (which is tiny), but the access time is extremely quick.
- When a program needs to work on a piece of data, it will bring it into a register first.

x86 Organization: Registers

- x86-64 contains 16 general purpose registers.
- They are identified by a short name such as (%rax, %rsi, %r8, etc.)
- In the interest of being able to access each register at a finer grain, there are multiple ways of accessing the data within each register.
- Ex. %rax refers to the full 64-bits stored in the register while %eax refers to the lower 32-bits.

x86 Organization: Registers

- `_h`: upper 8-bits of lower 16-bits
- `_l`: lower 8-bits
- `_x`: lower 16-bits.
- `e_x`: lower 32-bits (e stands for 'extended').
- `r_x`: full 64-bit register (r stands for...? 'rextended'? I don't know, they can't all be winners)

General Purpose Registers (A, B, C and D)

64	56	48	40	32	24	16	8
R?X							
				E?X			
						?X	
						?H	?L

Register Organization

- GP registers: %rax, %rbx, %rcx, %rdx, %r8~%r15
 - %rax: Accumulator, return value
 - %rbx: Base index (for use with arrays)
 - %rcx: Counter (for use with loops and strings), integer argument 0
 - %rdx: Data, integer argument 1
 - %r8: integer argument 2, %r9: integer argument 3
- Index registers:
 - %rsi: *Source index* for [string](#) operations.
 - %rdi: *Destination index* for string operations.
- Pointer registers:
 - %rsp: Stack pointer for top address of the stack.
 - %rbp: Stack base pointer for holding the address of the current [stack frame](#).
- Instruction pointer register:
 - %rip: Instruction pointer. Holds the [program counter](#), the current instruction address.

x86 Organization: Registers

- Registers are very simple containers that store a bit configuration and nothing more.
- If you know that a 64-bit signed long is stored in a register, there is nothing about this register that indicates that the original intention was for this value to be signed.
- The compiler will simply compile machine code that treats the bit vector in the register as if it were a signed value.

x86 Assembly

- What does this (GNU/GAS/AT&T) syntax look like?
- [operation] [source] [destination]
- ex.
- `movq %rax 4(%rbx)`
- `addq %rbx %rcx`

x86 Assembly: Ops

- The `mov_` family: move data from the source to the destination. The suffix determines how much data to move.
 - `movb` : move a byte
 - `movw` : move a word (16 bits (?))
 - `movl` : move a long/double word (32 bits)
 - `movq` : move a quad word (64 bits)
- `movq %rax, %rbx`
- `movq %rax, (%rbx)`

x86 Assembly: Ops

- A comment regarding “word” size.
- A “word” is just a label of convenience that is used refer to contiguous bytes of memory of a common size.
- On a 32-bit machine, a word is 4 bytes. On a 64-bit machine, a word is 8 bytes.
- But back when registers were only 16-bits and x86 was being developed, “words” were 16-bits.
- Dark days...

x86 Assembly: Addressing Modes

- Consider:
 - `movq %rax, (%rbx) <--- ?`
- The parentheses indicate a memory operation.
- That is, the source and destination operands are able to refer to values that are located in memory, rather than just registers.
- The parentheses `()` means:
 - Treat the bit vector within as a memory address.
 - Go follow that address into memory and get the value at that address.

x86 Assembly: Addressing Modes

- Say `%rax = 0xFEEDABBA` and `%rbx = 0x80`.
- `movq %rax, %rbx`
 - Result: `%rax = 0xFEEDABBA`, `%rbx = 0xFEEDABBA`
- `movq %rax, (%rbx)`
 - Result: the value that is located in memory address `0x80` is set as `0xFEEDABBA`.
 - In a more C-like form, this is essentially:
 - `MEM[0x80] = 0xFEEDABBA`; or
 - `*(0x80) = 0xFEEDABBA`;

x86 Assembly: Addressing Modes

- The '\$' symbol prefix indicates an “immediate” which is constant number value.
- If %rax = 0xb1ab
- `movl $0xdea1, %rax`
 - Result: %rax = 0xdea1

x86 Assembly: Basic insns

- Other instructions:
- `add_src, dst` `# dst = dst + src`
- `sub_src, dst` `# dst = dst - src`
- `neg_dst` `#dst = -dst`
- `not_dst` `#dst = ~dst`
- `sar imm, dst` `# dst = dst >> imm` (shift arithmetic right)
- Etc... (pg. 192)

x86 Assembly: Advanced Addressing Modes

- IMM(R1, R2, S) : Scaled and displaced array access.
 - Intended usage:
 - R1 : Base array address
 - R2 : Index into array
 - S : Size of array data type in bytes
 - IMM : Displacement
- movq A(B, C, D), %rax
 - $\%rax = *(A + B + C * D)$
- See pg 181 for full list

x86 Assembly: Advanced Addressing Modes

- D(R1) : Base + displacement addressing
 - If %rax = 0x10.
 - `movq 8(%rax), %rbx`
 - Result: %rbx gets the value at memory address $0x10 + 8 = 0x18$, *not* the number 0x18
 - $\%rbx = *(\%rax + 8)$.

x86 Assembly: Advanced Addressing Modes

- $(,R2,S)$: Scaled indexing, s must be 1, 2, 4, or 8.
 - If $\%rax = 0x01$ and $S = 2$.
 - `movl (, %rax, 2), %rcx`
 - Result: The value at memory address $0x01 * 2$ is saved to $\%rcx$.
- $IMM(R1, R2, S)$: Scaled and displaced array access.
 - If $\%rax = 0x400$, $\%rbx = 2$, $S = 2$, and $D = 20$.
 - `movl 0x20(%rax, %rbx, 2), %rcx`
 - Result: The value at memory address $0x400 + 0x2*2 + 0x20$ is placed in $\%rcx$.

x86 Assembly: lea_

- “lea_” or the equally-as-confusing full name “load effective address” is an unusual instruction that slightly violates the normal behavior.
- `movq (%rax), %rbx` => The value at memory address contained by %rax is moved to %rbx.
- `leaq (%rax), %rbx` => The value in %rax itself is moved to %rbx.
- What's the point of this?

x86 Assembly: lea_

- Consider:
- `movq 4(%rax,%rbx, 2), %rcx.`
 - This means $\%rcx = \text{MEM}[\%rax + \%rbx * 2 + 4]$
- `leaq 4(%rax,%rbx, 2), %rcx.`
 - This means $\%rcx = \%rax + \%rbx * 2 + 4$
- It's just a shortcut. This is a faster way to do address calculation or just to do arithmetic.

x86 Assembly nuances: mov_

- It is possible to move from a smaller container to a larger container.
- Assume that %dh = 0xCD, %eax = 0x98765432
- movb %dh, %eax
 - What's the result?

x86 Assembly nuances: mov_

- It is possible to move from a smaller container to a larger container.
- Assume that %dh = 0xCD, %eax = 0x98765432
- movb %dh, %eax
 - What's the result? It's not allowed (suffix mismatch). The destination has a 32-bit length while the movb expects only to move a byte.
- movb %dh, %al
 - Result: %eax = 0x987654CD
- However, you may want to extend the value in the source (ie %eax = 0xCD).

x86 Assembly nuances: mov_

- `movsXY` : move and sign extend from size X to size Y.
 - Ex: `movsbl` : move a byte from the source and sign extend it to a long (4 bytes)
- `movzXY` : move and zero extend from size X to size Y.
 - Ex: `movzbw` : move a byte from the source and zero extend it to a word (2 bytes)

x86 Assembly nuances: mov_

- %dh = 0xCD, %eax = 0x98765432
- movsbl %dh, %eax
 - Result: %eax = 0xFFFFFFFFCD
- movzbl %dh, %eax
 - Result: %eax = 0x000000CD
- movzbw %dh, %eax?
 - Result: Not allowed. Sign extending to w (16-bits) but to %eax (32-bit container).

x86 Assembly nuances: mov_

- As a final note about mov, the size of the prefix must match the operands. You cannot have:
 - `movl %ax, (%esp)` // Can't move a 32-bit quantity from a 16-bit register
 - `movl %eax, %dx` // Can't move a 32-bit quantity into a 16-bit register.

x86 Assembly nuances: mov_

- Additionally memory references match all sizes:
 - `movb %al, (%rbx)`
 - `movw %ax, (%rbx)`
 - `movq %rax, (%rbx)`
 - All allowed. The data will be moved to memory starting at that address based on the data type size
- However:
 - `movb %al, (%ebx)`
 - Is not meaningful on x86-64 because the `%ebx` register is only the lower 32-bits while addresses are 64-bits.